Problem Statement: Abalone Age Prediction

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Load the dataset

data.head()

```
data=pd.read_csv("abalone.csv")
```

```
Sex Length Diameter Height Whole weight Shucked weight Viscera weight
\
0
   Μ
        0.455
                  0.365
                          0.095
                                       0.5140
                                                       0.2245
                                                                       0.1010
1
   Μ
       0.350
                  0.265
                          0.090
                                       0.2255
                                                       0.0995
                                                                       0.0485
2
                  0.420
       0.530
                          0.135
                                       0.6770
                                                       0.2565
                                                                       0.1415
3
       0.440
                  0.365
                          0.125
                                       0.5160
                                                       0.2155
                                                                       0.1140
   Μ
4
   Ι
       0.330
                  0.255
                          0.080
                                       0.2050
                                                       0.0895
                                                                       0.0395
```

```
Shell weight Rings
0
          0.150
                     15
1
          0.070
                     7
2
                     9
          0.210
3
          0.155
                     10
4
          0.055
                      7
```

data.shape

(4177, 9)

data.size

37593

data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):

Ducu	COTAMINS (COCAT	J COTAMITS).	
#	Column	Non-Null Count	Dtype
0	Sex	4177 non-null	object
1	Length	4177 non-null	float64
2	Diameter	4177 non-null	float64
3	Height	4177 non-null	float64
4	Whole weight	4177 non-null	float64
5	Shucked weight	4177 non-null	float64
6	Viscera weight	4177 non-null	float64

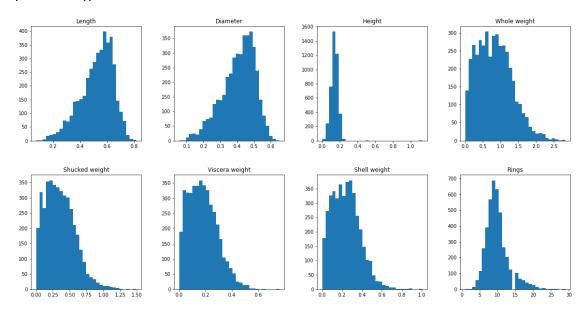
7 Shell weight 4177 non-null float64 8 Rings 4177 non-null int64 dtypes: float64(7), int64(1), object(1)

memory usage: 293.8+ KB

Data Analysis & Visualization

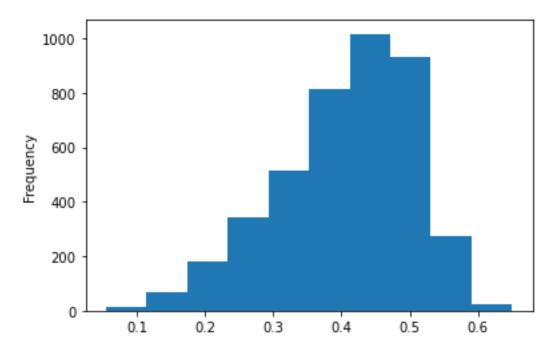
Univariate Analysis

data.hist(figsize=(20,10), grid=False, layout=(2,4), bins=30)
plt.show()

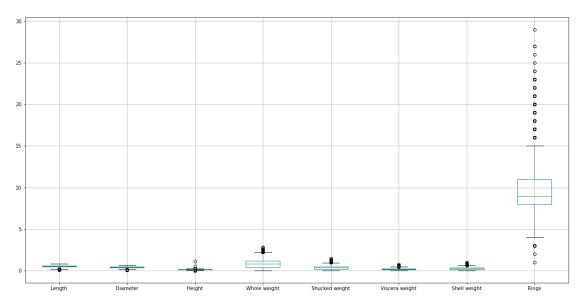


data["Diameter"].plot(kind='hist')

<AxesSubplot:ylabel='Frequency'>



data.boxplot(figsize=(20,10))
plt.show()




Bivariate Analysis data.head()

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
\							
0	Μ	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	Μ	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	Μ	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395

```
Shell weight Rings
0
             0.150
                           15
             0.070
1
                            7
2
             0.210
                            9
3
             0.155
                           10
4
             0.055
                            7
fig, axes = plt.subplots(4,2, figsize=(15,15))
axes = axes.flatten()
for i in range(1,len(data.columns)-1):
     sns.scatterplot(x=data.iloc[:,i], y=data['Rings'], ax=axes[i])
plt.show()
                                                   25
  0.8
                                                   20
  0.6
                                                 Sing 15
                                                   10
  0.2
 0.0
           0.2
                    0.4
                            0.6
                                    0.8
                                                                                0.6
                                                                                     0.7
  30
                                                   25
  25
  20
                                                   20
Rings
15
                                                 Sing 15
  10
                    0.3 0
Diameter
                                       0.6
                                                                                       1.0
                                                                        Height
  30
                                                   30
  20
Sings
15
                                                 Sing 15
  10
                                                   10
                                      2.5
                                                                                    1.2
                     Whole weight
                                                                     Shucked weight
  25
                                                   25
  20
                                                   20
Rings
15
                                                 Si 15
  10
                    0.3 0.4
Viscera weight
plt.figure(figsize=(10,5))
sns.boxenplot(y=data['Rings'], x=data['Sex'])
plt.grid()
```

plt.show()
data.groupby('Sex')['Rings'].describe()

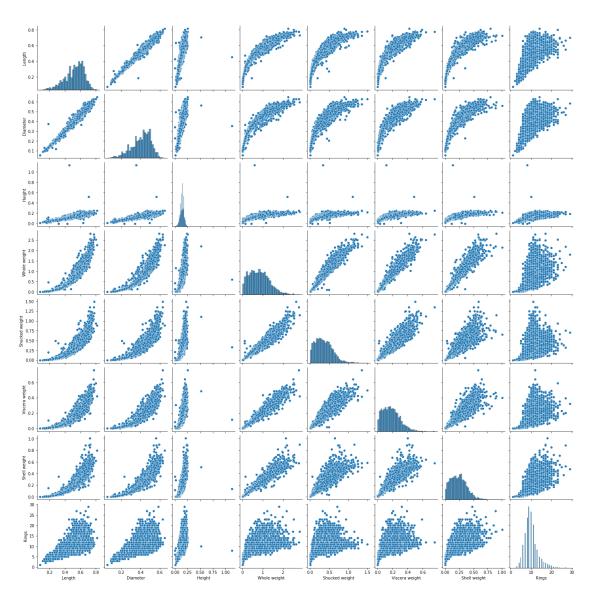


	count	mean	std	min	25%	50%	75%	max
Sex								
F	1307.0	11.129304	3.104256	5.0	9.0	10.0	12.0	29.0
I	1342.0	7.890462	2.511554	1.0	6.0	8.0	9.0	21.0
М	1528.0	10.705497	3.026349	3.0	9.0	10.0	12.0	27.0

Multivariate Analysis

sns.pairplot(data)

<seaborn.axisgrid.PairGrid at 0x161343cad00>



Descriptive statistics

data.describe()

	Length	Diameter	Height	Whole weight	Shucked weight	\
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	
mean	0.523992	0.407881	0.139516	0.828742	0.359367	
std	0.120093	0.099240	0.041827	0.490389	0.221963	
min	0.075000	0.055000	0.000000	0.002000	0.001000	
25%	0.450000	0.350000	0.115000	0.441500	0.186000	
50%	0.545000	0.425000	0.140000	0.799500	0.336000	
75%	0.615000	0.480000	0.165000	1.153000	0.502000	
max	0.815000	0.650000	1.130000	2.825500	1.488000	

Viscera weight Shell weight Rings count 4177.000000 4177.000000 4177.000000 mean 0.180594 0.238831 9.933684

std	0.109614	0.139203	3.224169
min	0.000500	0.001500	1.000000
25%	0.093500	0.130000	8.000000
50%	0.171000	0.234000	9.000000
75%	0.253000	0.329000	11.000000
max	0.760000	1.005000	29.000000

Handle The Missing values

data.isnull().any()

Sex False False Length False Diameter Height False Whole weight False Shucked weight False Viscera weight False Shell weight False False Rings

dtype: bool

data.isnull().sum()

Sex 0 Length 0 Diameter 0 Height 0 Whole weight 0 Shucked weight 0 Viscera weight 0 Shell weight 0 Rings 0 dtype: int64

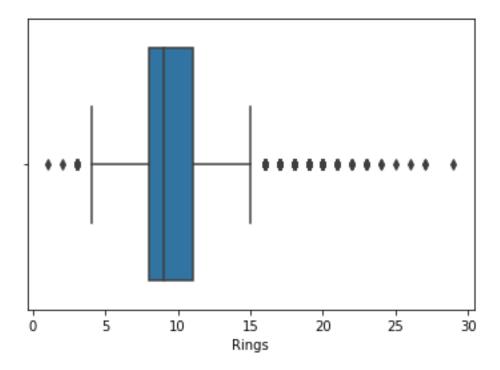
Find the outliers

data.skew()

Length -0.639873 Diameter -0.609198 Height 3.128817 Whole weight 0.530959 Shucked weight 0.719098 Viscera weight 0.591852 Shell weight 0.620927 1.114102 Rings

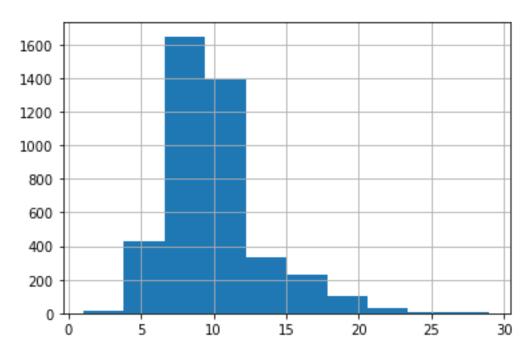
dtype: float64

sns.boxplot(x=data['Rings'],data=data)
<AxesSubplot:xlabel='Rings'>



data['Rings'].hist()

<AxesSubplot:>



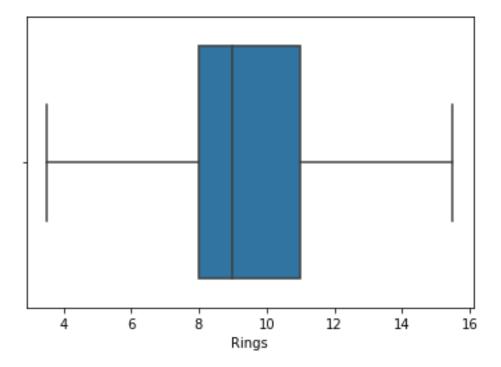
print('skewness value of Age: ',data['Rings'].skew())

skewness value of Age: 1.114101898355677

Flooring And Capping

```
Q1 = data['Rings'].quantile(0.25)
Q3 = data['Rings'].quantile(0.75)
IQR = Q3 - Q1
whisker_width = 1.5
lower_whisker = Q1 -(whisker_width*IQR)
upper_whisker = Q3 +(whisker_width*IQR)
data['Rings']=np.where(data['Rings']>upper_whisker,upper_whisker,np.where(dat
a['Rings']<lower_whisker,lower_whisker,data['Rings']))</pre>
sns.boxplot(x=data['Rings'],data=data)
```

<AxesSubplot:xlabel='Rings'>



Categorical encoding

data.info()

<class 'pandas.core.frame.DataFrame'> RangeIndex: 4177 entries, 0 to 4176 Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Sex	4177 non-null	object
1	Length	4177 non-null	float64
2	Diameter	4177 non-null	float64
3	Height	4177 non-null	float64
4	Whole weight	4177 non-null	float64
5	Shucked weight	4177 non-null	float64
6	Viscera weight	4177 non-null	float64
7	Shell weight	4177 non-null	float64

```
4177 non-null
     Rings
                                     float64
dtypes: float64(8), object(1)
memory usage: 293.8+ KB
#Label Encoding
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
data['Sex']=le.fit_transform(data['Sex'])
data.head()
   Sex Length Diameter Height Whole weight Shucked weight
     2
        0.455
                   0.365
                           0.095
                                        0.5140
                                                        0.2245
1
     2
        0.350
                   0.265
                           0.090
                                        0.2255
                                                        0.0995
2
    0 0.530
                   0.420
                           0.135
                                        0.6770
                                                        0.2565
3
    2 0.440
                   0.365
                           0.125
                                                        0.2155
                                        0.5160
4
    1
        0.330
                   0.255
                           0.080
                                        0.2050
                                                        0.0895
   Viscera weight Shell weight
                                 Rings
0
          0.1010
                          0.150
                                  15.0
1
           0.0485
                          0.070
                                   7.0
           0.1415
2
                          0.210
                                   9.0
3
                          0.155
                                  10.0
          0.1140
4
           0.0395
                          0.055
                                  7.0
data["Sex"].unique()
array([2, 0, 1])
Split the data into training and testing
data.head(5)
   Sex Length Diameter
                         Height
                                 Whole weight Shucked weight
0
     2
       0.455
                   0.365
                           0.095
                                        0.5140
                                                        0.2245
     2
        0.350
                   0.265
                           0.090
                                        0.2255
                                                        0.0995
1
2
    0 0.530
                   0.420
                           0.135
                                        0.6770
                                                        0.2565
3
    2
        0.440
                   0.365
                           0.125
                                        0.5160
                                                        0.2155
4
        0.330
                   0.255
                           0.080
                                        0.2050
                                                        0.0895
   Viscera weight Shell weight Rings
0
           0.1010
                          0.150
                                  15.0
1
           0.0485
                          0.070
                                   7.0
2
           0.1415
                          0.210
                                   9.0
           0.1140
3
                          0.155
                                  10.0
```

0.055

7.0

0.0395

```
X = data.iloc[:, 0:7]
Y = data.iloc[:,-1]
Χ
                    Diameter
                               Height
                                       Whole weight Shucked weight
      Sex
           Length
0
        2
            0.455
                       0.365
                                0.095
                                              0.5140
                                                               0.2245
                       0.265
1
        2
            0.350
                                0.090
                                              0.2255
                                                               0.0995
2
            0.530
                       0.420
        0
                                0.135
                                              0.6770
                                                               0.2565
3
        2
             0.440
                       0.365
                                0.125
                                              0.5160
                                                               0.2155
4
        1
             0.330
                       0.255
                                0.080
                                              0.2050
                                                               0.0895
               . . .
                          . . .
                                  . . .
4172
        0
            0.565
                       0.450
                                0.165
                                              0.8870
                                                               0.3700
                       0.440
4173
        2
            0.590
                                0.135
                                              0.9660
                                                               0.4390
4174
        2
            0.600
                       0.475
                                0.205
                                              1.1760
                                                               0.5255
4175
        0
             0.625
                       0.485
                                0.150
                                              1.0945
                                                               0.5310
4176
        2
             0.710
                       0.555
                                0.195
                                              1.9485
                                                               0.9455
      Viscera weight
0
               0.1010
1
               0.0485
2
               0.1415
3
               0.1140
4
               0.0395
               0.2390
4172
4173
               0.2145
4174
               0.2875
4175
               0.2610
4176
               0.3765
[4177 rows x 7 columns]
Υ
0
        15.0
1
         7.0
2
         9.0
3
        10.0
4
         7.0
        . . .
4172
        11.0
4173
        10.0
4174
         9.0
4175
        10.0
4176
        12.0
Name: Rings, Length: 4177, dtype: float64
y.shape
(4177, 0)
```

Scale the independent variables

```
from sklearn.preprocessing import StandardScaler
ss = StandardScaler()
X_scaled = ss.fit_transform(X)
```

Split the data into training and testing

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X_scaled, Y, test_size =
0.3, random_state = 1)
```

Build the Model

```
Training the Model Testing the Model
import csv
with open("abalone.csv") as csv_file:
    csv_reader = csv.reader(csv_file)
    data = pd.DataFrame([csv_reader], index = None)
for val in list(data[1]):
    print(val)

['M', '0.455', '0.365', '0.095', '0.514', '0.2245', '0.101', '0.15', '15']
```

Training and Testing Module

1.Linear Regression 2.Ridge 3.Decision Tree Regression 4.KNeighborsRegressor

```
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor

from sklearn.metrics import mean_squared_error, r2_score

1. Linear Regression
lr = LinearRegression()
lr.fit(x_train, y_train)
LinearRegression()
#Testing the model
lr_test_pred = lr.predict(x_test)
lr test_pred
```

#importing all the neccessary models and metrics

```
array([8.49722433, 7.64369059, 7.82520883, ..., 8.55677832, 9.02884473,
       5.96561877])
#measuring the performance
mse = mean_squared_error(y_test, lr_test_pred)
print('Mean Squared error of testing Set: %2f'%mse)
Mean Squared error of testing Set: 3.524602
p = r2_score(y_test, lr_test_pred)
print('R2 Score of testing set:%.2f'%p)
R2 Score of testing set:0.52
   Ridge
ridge mod = Ridge(alpha=0.01, normalize=True)
ridge_mod.fit(x_train, y_train)
ridge_mod.fit(x_test, y_test)
Ridge(alpha=0.01, normalize=True)
#Testing the model
ridge_model_pred = ridge_mod.predict(x_test)
ridge model pred
array([8.54031033, 8.48463396, 7.96838487, ..., 8.77493484, 9.03881023,
       5.83582085])
#Measuring the performance
acc = r2_score(y_test, ridge_model_pred)
print('Score of testing Set: %2f'%acc)
Score of testing Set: 0.523227
   Decision Tree Regression
dt = DecisionTreeRegressor()
dt.fit(x_train, y_train)
DecisionTreeRegressor()
#Testing the model
dt_test_pred = dt.predict(x_test)
dt_test_pred
array([12., 9., 10., ..., 7., 9., 4.])
#Measuring the Performance
dacc = mean_squared_error(y_test, dt_test_pred)
print('Mean Squared Error of testing Set: %2f'%dacc)
```

```
Mean Squared Error of testing Set: 6.126994
1. KNN Regression
knn = KNeighborsRegressor(n_neighbors = 4 )
knn.fit(x_train, y_train)
knn.fit(x_test, y_test)
KNeighborsRegressor(n_neighbors=4)
#Testing the Model
knn_test_pred = knn.predict(x_test)
knn_test_pred
array([ 8.75, 9.5, 10.5, ..., 8. , 7.5, 5. ])
#Measuring the Performance
kacc= r2_score(knn_test_pred,y_test)
print('Score of testing Set: %2f'%kacc)
Score of testing Set: 0.400555
kmse = mean_squared_error(knn_test_pred,y_test)
print('Score of testing Set: %2f'%kmse)
Score of testing Set: 2.602460
```